

```

import os
import cv2
import numpy as np
from sklearn.svm import SVC
from sklearn.metrics import classification_report, accuracy_score,
confusion_matrix
from sklearn.decomposition import PCA
from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Model
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from skimage.feature import local_binary_pattern, hog
import matplotlib.pyplot as plt
import seaborn as sns

# Paths
TRAIN_PATH = "/Users/lokeshtarinath/Downloads/archive (3)/train"
TEST_PATH = "/Users/lokeshtarinath/Downloads/archive (3)/test"
IMG_SIZE = (48, 48)

# Utility: Load and preprocess images
def load_data(base_path):
    data, labels = [], []
    for class_label, class_name in enumerate(os.listdir(base_path)):
        folder_path = os.path.join(base_path, class_name)
        if os.path.isdir(folder_path):
            for file_name in os.listdir(folder_path):
                img_path = os.path.join(folder_path, file_name)
                img = cv2.imread(img_path, cv2.IMREAD_GRAYSCALE)
                if img is not None:
                    img_resized = cv2.resize(img, IMG_SIZE)
                    data.append(img_resized)
                    labels.append(class_label)
    return np.array(data, dtype="float32"), np.array(labels)

# Load datasets
X_train, y_train = load_data(TRAIN_PATH)
X_test, y_test = load_data(TEST_PATH)

# Normalize pixel values
X_train /= 255.0
X_test /= 255.0

# Data augmentation
datagen = ImageDataGenerator(
    rotation_range=10,
    width_shift_range=0.1,
    height_shift_range=0.1,
    brightness_range=[0.8, 1.2],
    horizontal_flip=True
)

```

```

datagen.fit(X_train[..., np.newaxis])

# Baseline SVM model on raw images
svm_baseline = SVC(kernel='rbf', C=1, gamma=0.1)
svm_baseline.fit(X_train.reshape(len(X_train), -1), y_train)
y_pred_baseline = svm_baseline.predict(X_test.reshape(len(X_test), -1))

print("Baseline Model Metrics (Raw Images):")
print("Accuracy:", accuracy_score(y_test, y_pred_baseline))
print("Classification Report:\n", classification_report(y_test, y_pred_baseline))

# Visualize confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(confusion_matrix(y_test, y_pred_baseline), annot=True,
            fmt='d', cmap="Blues")
plt.title("Baseline Model Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

# Feature extraction methods
def extract_lbp_features(images):
    return np.array([
        np.histogram(local_binary_pattern(img, P=8, R=1,
            method="uniform").ravel(), bins=np.arange(0, 59), density=True)[0]
        for img in images
    ])

def extract_hog_features(images):
    return np.array([
        hog(img, orientations=9, pixels_per_cell=(8, 8),
            cells_per_block=(2, 2), visualize=False)
        for img in images
    ])

def extract_cnn_features(images):
    base_model = VGG16(weights="imagenet", include_top=False,
        input_shape=(48, 48, 3))
    feature_model = Model(inputs=base_model.input,
        outputs=base_model.layers[-1].output)
    preprocessed = np.array([cv2.cvtColor((img *
        255).astype(np.uint8), cv2.COLOR_GRAY2RGB) for img in images])
    cnn_features = feature_model.predict(preprocessed / 255.0,
        verbose=0)
    return cnn_features.reshape(len(images), -1)

# Extract features
lbp_train, lbp_test = extract_lbp_features(X_train),

```

```

extract_lbp_features(X_test)
hog_train, hog_test = extract_hog_features(X_train),
extract_hog_features(X_test)
cnn_train, cnn_test = extract_cnn_features(X_train),
extract_cnn_features(X_test)

# Dimensionality reduction with PCA
def apply_pca(features, n_components=50):
    pca = PCA(n_components=n_components)
    reduced_train = pca.fit_transform(features)
    return reduced_train, pca

lbp_train_pca, pca_lbp = apply_pca(lbp_train)
lbp_test_pca = pca_lbp.transform(lbp_test)

hog_train_pca, pca_hog = apply_pca(hog_train)
hog_test_pca = pca_hog.transform(hog_test)

cnn_train_pca, pca_cnn = apply_pca(cnn_train)
cnn_test_pca = pca_cnn.transform(cnn_test)

# Combine features
combined_train = np.hstack([lbp_train_pca, hog_train_pca,
cnn_train_pca])
combined_test = np.hstack([lbp_test_pca, hog_test_pca, cnn_test_pca])

# SVM with combined features
svm_combined = SVC(kernel='rbf', C=1, gamma=0.1)
svm_combined.fit(combined_train, y_train)
y_pred_combined = svm_combined.predict(combined_test)

print("Final Model Metrics (Combined Features):")
print("Accuracy:", accuracy_score(y_test, y_pred_combined))
print("Classification Report:\n", classification_report(y_test,
y_pred_combined))

# Visualize confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(confusion_matrix(y_test, y_pred_combined), annot=True,
fmt='d', cmap="Greens")
plt.title("Combined Features Model Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

# Feature visualization
plt.figure()
plt.hist(local_binary_pattern(X_train[0], P=8, R=1,
method="uniform").ravel(), bins=np.arange(0, 59))

```

```
plt.title("LBP Histogram (Sample Image)")
plt.show()
```


```
/Users/lokeshharinath/Library/Python/3.9/lib/python/site-packages/
urllib3/__init__.py:35: NotOpenSSLWarning: urllib3 v2 only supports
OpenSSL 1.1.1+, currently the 'ssl' module is compiled with 'LibreSSL
2.8.3'. See: https://github.com/urllib3/urllib3/issues/3020
  warnings.warn(
```

Baseline Model Metrics (Raw Images):

Accuracy: 0.34215658957927

Classification Report:

	precision	recall	f1-score	support
0	0.28	0.99	0.43	1774
1	0.69	0.06	0.12	1247
2	0.89	0.13	0.23	1024
3	0.98	0.31	0.47	831
4	0.75	0.09	0.16	1233
5	0.89	0.09	0.16	958
6	1.00	0.28	0.44	111
accuracy			0.34	7178
macro avg	0.78	0.28	0.29	7178
weighted avg	0.69	0.34	0.27	7178



```
/Users/lokeshharinath/Library/Python/3.9/lib/python/site-packages/skimage/feature/texture.py:360: UserWarning: Applying `local_binary_pattern` to floating-point images may give unexpected results when small numerical differences between adjacent pixels are present. It is recommended to use this function with images of integer dtype.
```

```
warnings.warn(
```


Final Model Metrics (Combined Features):

Accuracy: 0.5483421565895793

Classification Report:

	precision	recall	f1-score	support
0	0.56	0.79	0.66	1774
1	0.44	0.48	0.46	1247
2	0.57	0.38	0.45	1024
3	0.78	0.66	0.71	831
4	0.48	0.51	0.50	1233
5	0.54	0.35	0.43	958

	6	1.00	0.35	0.52	111
accuracy				0.55	7178
macro avg		0.63	0.50	0.53	7178
weighted avg		0.56	0.55	0.54	7178



```

/Users/lokeshharinath/Library/Python/3.9/lib/python/site-packages/
skimage/feature/texture.py:360: UserWarning: Applying
`local_binary_pattern` to floating-point images may give unexpected
results when small numerical differences between adjacent pixels are
present. It is recommended to use this function with images of integer
dtype.
  warnings.warn(

```

LBP Histogram (Sample Image)

